



WHERE ARE WE NOW?

W34: Introduction, DBMSs and Relational Databases

W35: Developing Database Systems

W36: SQL –Part I

W37: SQL –Part II and Relational Algebra

W38: Data Modelling

W39: Data Modelling

W40: Database Design

W41: Normalisation and Stored Procedures

W42: XML and Web Technology

W43: Processing XML Data

W44: XML Validation

W45: Beyond relational databases and XML

W46: File Organisations and Indexes

W47: Database Security and Administration

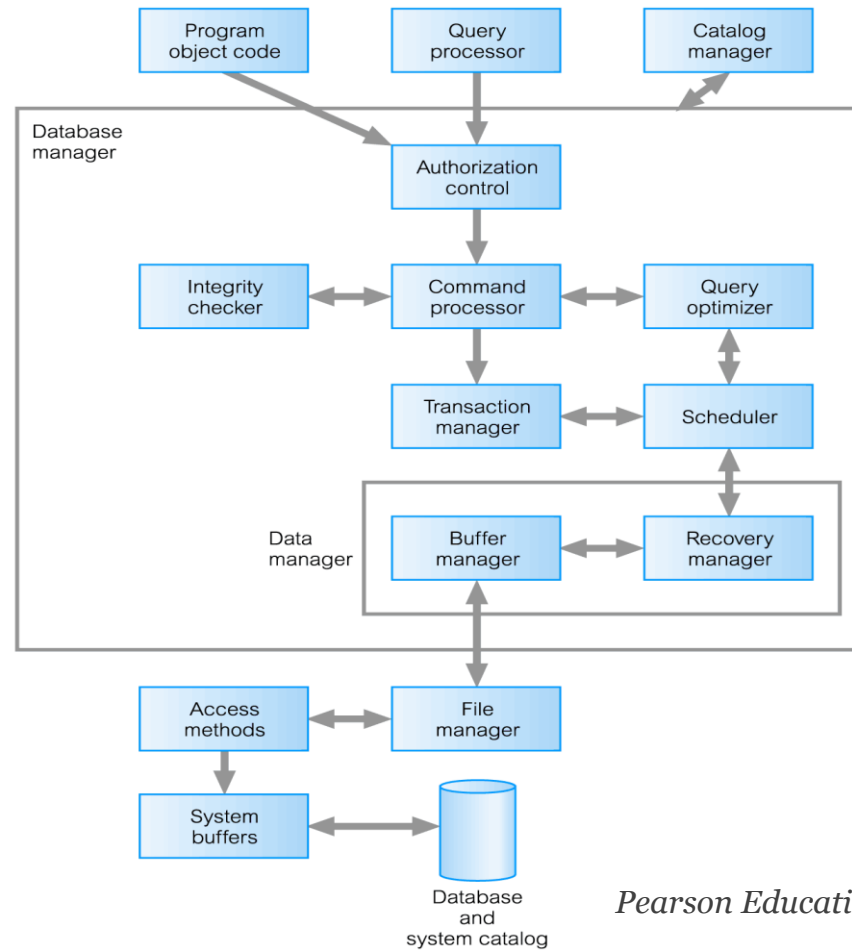
W48: Transaction Processing and Wrap-up



CONTENT

- *Database file organisation and access methods:*
 - » Introduction
 - » Heaps
 - » Ordered files
 - » Hash files
- *Indexes and index file organisation:*
 - » Primary and secondary indexes
 - » B+-trees
 - » Bitmap indexes
- *Managing Indexes in SQL*

DATABASE MANAGER



Pearson Education © 2009

STORAGE MEDIA (1)

- Primary vs secondary storage:
 - *Primary storage:*
 - » Fast random access (ns)
 - » Volatile
 - » Expensive
 - *Secondary storage:*
 - » “Fast” (ms) - six orders of magnitude slower than main memory
 - » Persistent (non-volatile) data storage
 - » “Inexpensive” - two orders of magnitude cheaper than main memory

STORAGE MEDIA(2)

- DBMSs store data on disks:
 - *Although some real-time database systems rely on in-memory databases*
- Disk characteristics:
 - *Unit of read/write operations:*
 - » a logical block/page storing several rows
 - *Access time:*
 - » seek time + rotation time + transfer time
 - *Sequential I/O much faster than random I/O*

STORAGE MEDIA(3)

- Methods to improve main memory-disk data transfer:
 - *General approaches:*
 - » Improve the disk technology
 - » Use faster disks (more RPMs)
 - » Parallelization (RAID) + some redundancy
 - » Other (OS based improvements): disk scheduling (elevator algorithm, batch writes, etc)
 - *DBMS approaches:*
 - » Good file organization
 - » Avoid unnecessary reads from disk
 - » Buffer management: go to buffer instead of disk

INTRODUCTION TO FILE ORGANIZATION (1)

- Basics:
 - *A database holds a collection of files, file holds a collection of records, record (tuple) is a collection of fields (attributes)*
 - *Some database systems (e.g., MySQL and Oracle) allow files to be grouped into tablespaces*
 - *Files are stored on disks*
- Two important issues:
 - *Representation of each record*
 - *Grouping/Ordering of records and storage in blocks*

INTRODUCTION TO FILE ORGANIZATION (2)


- Goals and considerations:
 - *Compactness*
 - *Overhead of insertion/deletion*
 - *Retrieval speed:*
 - » sometime we prefer to bring more tuples than necessary into main memory and use CPU to filter out the unnecessary ones!

DATABASE RECORD REPRESENTATION (1)

- Fixed-size records:
 - *Store record i starting from byte $n * (i - 1)$, where n is the size of each record.*
 - *Might not allow records to cross block boundaries*
 - *Record access is simple*
 - *Example:*
 - » *Ordre (kundeid int, dato date, varenummer char(16), antall int)*
 - » *Size n in bytes: $4 + 3 + 16 + 4 = 27$*
 - » *A 4KB block can hold $\lfloor 4096/27 \rfloor = 151$ records*
 - ♦ *Leaving 19 unused bytes per block ($4096 - 151*27$)*


DATABASE RECORD REPRESENTATION (2)

- Variable-size records:
 - *There are alternative representations involving end-of-field characters and/or pointers*
 - *Common approach:*
 - » 1- or 2-byte pointers in the header point to the start/end of each field
 - » Example:
 - ♦ *Kunde(kundeld int, fornavn varchar(128),
etternavn varchar(128), epost varchar(128),
postnummer unsigned zerofilled smallint)*



DATABASE RECORD REPRESENTATION (3)

- Long attribute values (BLOB, CLOB, TEXT, ...) may be stored in a separate linked structure



Hammer & Schneider 2001



A NUT TO CRACK

- Another way to store records is to store them column-wise:
 - *Each column in different files*
 - *Advantages?*
 - *Disadvantages?*
- Column-oriented databases are being used for:
 - *Data warehousing*
 - *Data mining*
 - ...

HEAPS

- Simplest type
- Records stored in the order in which they were inserted
- New records added at the end of the file
- Suitable when:
 - *bulk loading*
 - *when typical access is a scan through all records*
- Records are marked as deleted when deleted:
 - *Record space is typically not reused*
 - *A periodic reorganisation may be required*

ORDERED FILES

- Records sorted on some key attribute(s):
 - *Guarantees record uniqueness if sorted on the primary key*
- New records added in correct position:
 - *May require a major move of records to create space*
- Suitable when:
 - *records must be returned in some order*
 - *a "range" of records needs to be retrieved*
- Records are reorganized when records are deleted

HASH FILES

- Records inserted in block based on hash value:
 - *By applying a hash function on some field(s)*
- New records added in correct block:
 - *Overflow mechanisms needed on collision:*
 - » open addressing
 - » unchained overflow
 - » chained overflow
 - » multiple hashing
- Suitable when:
 - *retrieving based on equality on the hash fields*
- Records may need to be reorganized when records are deleted



A NUT TO CRACK

- Assume that we need to do a complete sequential scan through all database records
- How will a hash file perform when compared to a heap file:
 - *Better*
 - *Worse*
 - *Similarly*



CONTENT

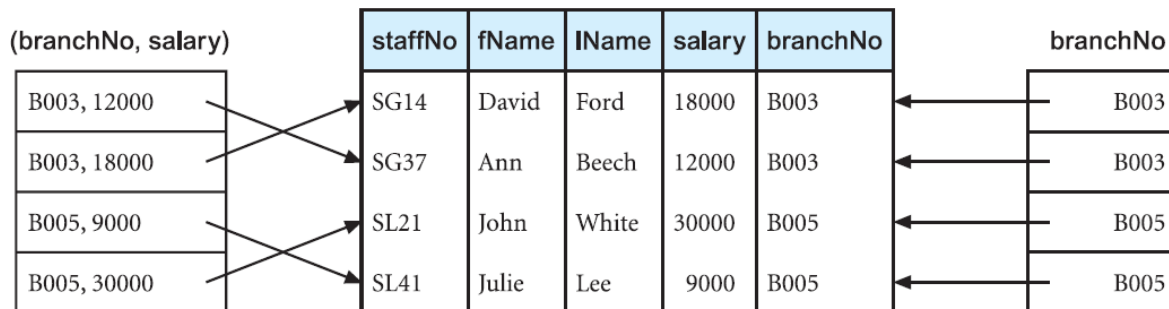
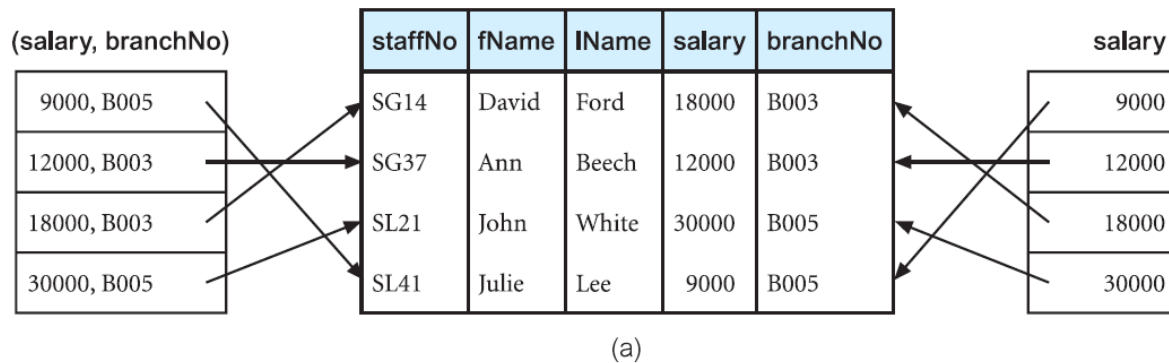
- *Database file organisation and access methods:*
 - » Introduction
 - » Heaps
 - » Ordered files
 - » Hash files
- *Indexes and index file organisation:*
 - » Primary and secondary indexes
 - » B+-trees
 - » Bitmap indexes
- *Managing Indexes in SQL*

INDEXES

- Primary indexes:
 - *The data file is sequentially ordered by the primary key*
 - *Guarantees uniqueness*
 - *Can be only one per file*
- Secondary indexes:
 - *An index defined on a non-ordering field of the data file*
 - *Does not have to be unique*
 - » But may to enforce uniqueness
 - *Can be many per file*
 - *Reduces access time significantly for large tables*

INDEXED SEQUENTIAL FILES (ISAM)

- Data files sorted on primary key
- Separate index or indexes for random access




Pearson Education © 2009

B+-TREES (1)

- Trees are known to provide efficient search structures
- In databases access times depend on number of disk accesses:
 - *Advantageous to have "bushy", shallow trees:*
 - » Many indexes stored in one block
 - » The tree as shallow as possible
- The B+-tree
 - *A **B**alanced tree*

B+-TREES (2)



Pearson Education © 2009

B+-TREES (3)

- B+-tree rules:
 - *The root must have at least two children, unless being a leaf*
 - *Each "interior" node has between $n/2$ and n pointers and children*
 - *The tree is always balanced*
 - *Leaf nodes are linked in order of key values*

B+-TREES (4)

- The depth of a B+-tree:
 - *Assume that the field to be indexed is char(8), that pointers are 4 bytes large, and that the block size is 4096 B.*
 - *Size of a block containing n indexes:*
 - » $n \cdot (4 + 8) + 4$
 - *Number of indexes per block*
 - » $n = 1 + (4096 - 4) / (4 + 8) = 342$
 - *Number of records that can be indexed:*
 - » 1 level: 342
 - » 2 levels: $342 * 341 = 116,622$
 - » 3 levels: $342 * 342 * 341 = 39,884,724$
 - » ...
 - » h levels: $n^h - n^{h-1}$

BITMAP INDEXES

- Useful for attributes having sparse domains
- Example: Music categories of a CD:

id	title	artist	genre	creationYear
1	Believe - Deluxe Edition (m/DVD)	Justin Bieber	Pop	2012
2	Contakt	Madcon	Hip	2012
3	Live Viking Stadion 9. Juni 2012	Mods	Rck	2012
4	Living Things	Linkin Park	HRk	2012
5	Odyssey - In Studio & In Concert (3CD)	Terje Rypdal	Kls	2012
6	Sexual Harassment - Limited Edition	Turboneger	Pun	2012
7	Slipp mine fløyter fri	Skruk	Kor	1990

Hip	HRk	Kls	Kor	Pop	Pun	Rck
0	0	0	0	1	0	0
1	0	0	0	0	0	0
0	0	0	0	0	0	1
0	1	0	0	0	0	0
0	0	1	0	0	0	0
0	0	0	0	0	1	0
0	0	0	1	0	0	0



CONTENT

- *Database file organisation and access methods:*
 - » Introduction
 - » Heaps
 - » Ordered files
 - » Hash files
- *Indexes and index file organisation:*
 - » Primary and secondary indexes
 - » B+-trees
 - » Bitmap indexes
- *Managing Indexes in SQL*

INDEXES IN SQL

- Creation of indexes is **not** standard in SQL
 - *But all vendors offer a CREATE INDEX statement*
- Indexes may be unique but do not have to be
- MySQL statement syntax:
 - *CREATE [UNIQUE] INDEX index_name
ON tbl_name (index_col_name,...)*
 - *DROP INDEX index_name ON tbl_name*
- We will cover indexes in more detail on Wednesday



RESOURCES

- C&B 7.3.5, Appendix F